

Introduction to Parallel Programming and OpenMP

Ilya A. Kaliman

August 1, 2017

Outline


- Lecture on the basics of parallel programming and an introduction to OpenMP
- Hands-on OpenMP tutorial
- Work on your own project

What is parallel computing?

Serial computing



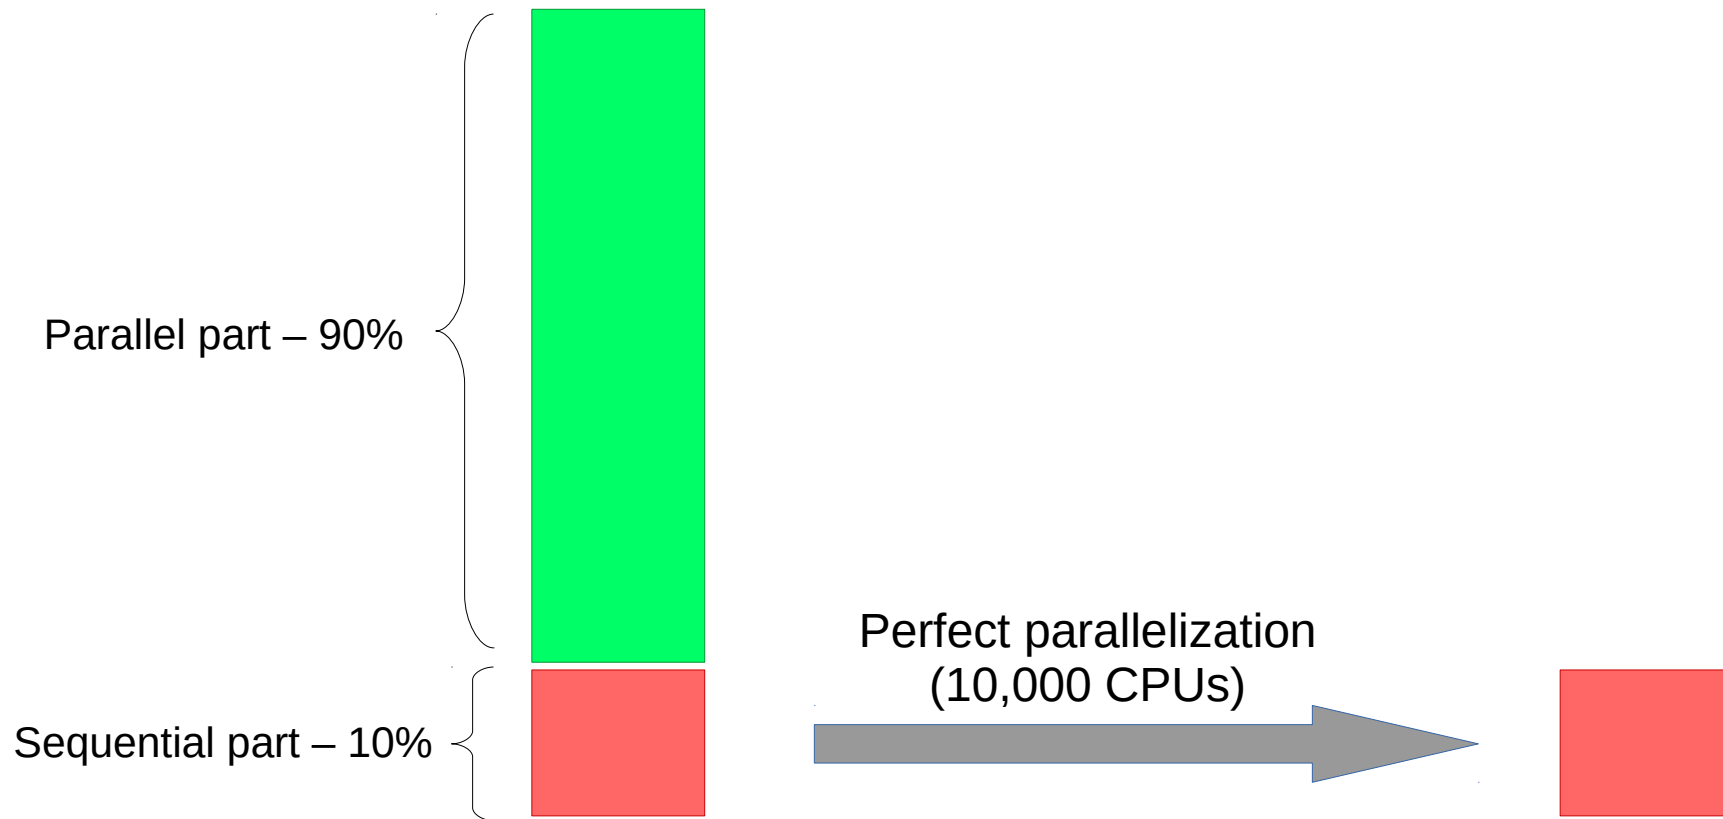
Parallel computing




Why parallelize?

- Solve problems in less time
 - Days → hours
- Solve larger problems
 - 100 atoms → 10,000 atoms
- Concurrently do several things
- Make more efficient use of hardware
 - All modern systems have multiple CPUs/cores
 - Performance of a single core does not grow


Parallelization limits: Amdahl's Law



Parallelization limits: Amdahl's Law



Parallelization limits: Amdahl's Law



$$t(P) = t_{seq} + \frac{t_{par}}{P}$$

$$Speedup(P) = \frac{t(1)}{t(P)}$$

t_{seq} – sequential portion

t_{par} – parallel portion

P – number of CPUs

Figure: wikipedia

Parallel programming challenges

- Dead/live locks
- Race conditions
- Resource oversubscription
- Algorithm/data structure design
- Complicated debugging
- ...

Simplest form of parallelism: bash

- `bash $: ./prog input1 > output1 &`
- `bash $: ./prog input2 > output2 &`
- `bash $: ./prog input3 > output3 &`
- `bash $: ./prog input4 > output4 &`

Parallel Programming Technologies

- Programming languages
 - Cilk, Chapel, Coarray Fortran, Charm++, ...
- Distributed systems
 - Message Passing Interface (MPI)
- Heterogeneous
 - OpenCL, CUDA
- Shared-memory
 - OpenMP

OpenMP

- OpenMP is an Application Programming Interface (API)
- Allows to develop **portable** and **scalable** applications on shared-memory platforms
- Consists of
 - Compiler directives (pragmas)
 - Runtime library (functions `omp_xxx_xxx`, `omp.h`)
 - Environment variables (`export OMP_NUM_THREADS=4`)
- Supports
 - C and C++
 - Fortran

Why OpenMP?

- Standard
- Well-established (supported by most C, C++, Fortran compilers)
- Relatively easy to use
- Straightforward path from sequential program to a parallel one
- Supported virtually everywhere (supercomputers)

Standards

- <http://www.openmp.org>
- Standards
 - OpenMP 4.5 (Nov 2015) – latest
 - OpenMP 3.1 (July 2011)
 -
 - OpenMP 1.0 (April 1999)
- Usually backwards compatible (adds new features without changing the behavior of older programs)

OpenMP: fork-join parallel model


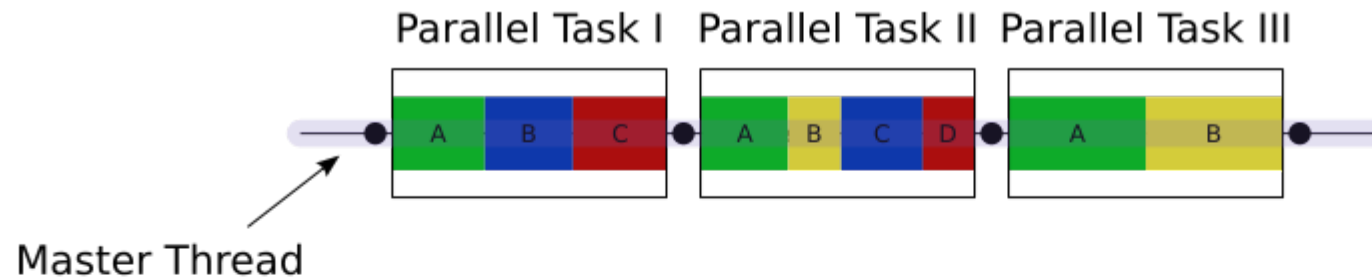


Figure: wikipedia

OpenMP Compiler Flags

- GNU (gcc, g++, gfortran)
 - **-fopenmp**
- Intel compiler (icc, icpc, ifort)
 - **-openmp** or **-qopenmp**
- Clang (clang, clang++)
 - **-fopenmp**
- Microsoft (cl.exe)
 - **/openmp**

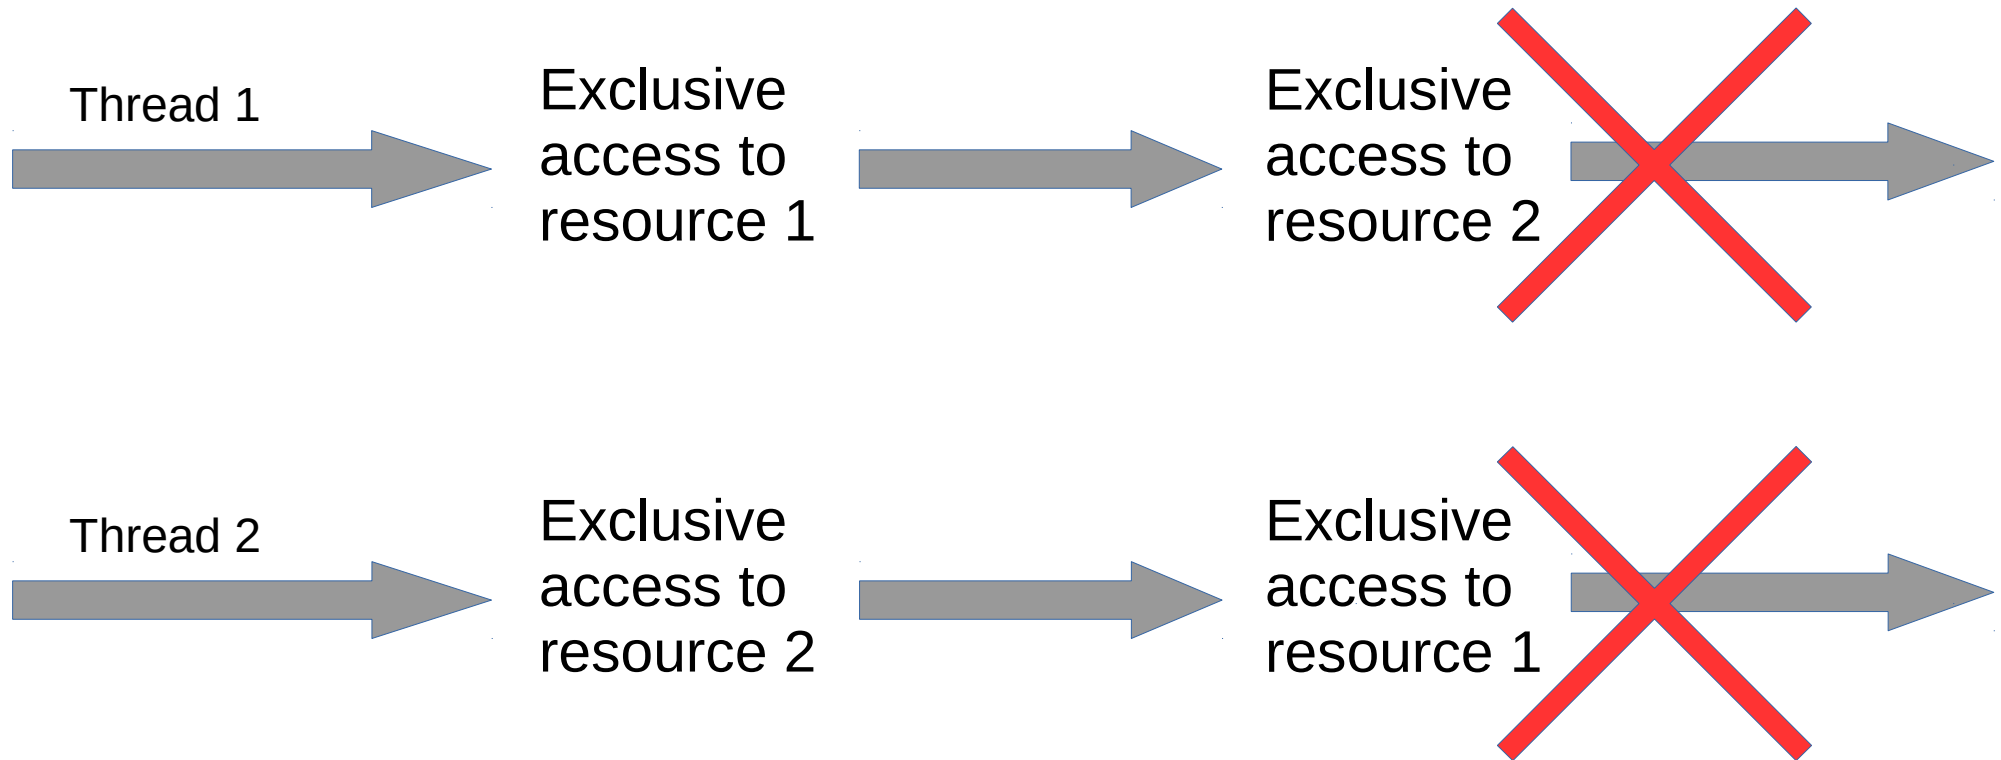
Data race

- We have in our code:
 - $x = x + 1$
- In reality:
 - Load x from memory to CPU register
 - Add 1
 - Store x to memory

Data race

- We have in our code:
 - $x = x + 1$
- In reality:
 - Load x from memory to CPU register
 - Add 1 ← **another thread may access memory here**
 - Store x to memory

Dead lock



GitHub repository with examples

```
git clone https://github.com/ilyak/openmp-tutorial
```